Jarvis

The control plane for AI execution

PROPOSE → APPROVE → EXECUTE → RECEIPT → TRACE

MARKET SHIFT

The shift

> **Agents are moving from generation to action**

- Send real outbound messages to customers
- Modify live systems and production data
- Trigger APIs, payments, and workflows

The question is not only “what did it say?” — it’s “what did it do?”

RISK

The problem

> **Execution is not trusted**

- Actions execute **without explicit human authorization**
- **Approval and execution collapse** into a single step
- **Logs ≠ proof** — no record of who actually authorized the action

That is a **governance and liability** problem: unbounded authority in production systems.

CATEGORY HEAT

Market validation

> **Enterprises are investing in agent infrastructure**

- Major platform and cloud investment in agent infrastructure
- **Registries and catalogs** — what exists, who owns it
- **Orchestration and platforms** — how agents are wired
- **Governance layers** — policy, visibility, reuse

Validates the category. It does not close the execution accountability gap.

The gap

> **Visibility ≠ control**

Platforms answer: *What can be discovered and shared?*

They do not answer:

> *What is allowed to happen — at the exact moment an agent takes action — under whose authority — with what proof?*

```
Agents & tools
      |
      ▼
Platforms (catalogs, orchestration)
      |
      ▼
  ┌──────────────┐
  │  EXECUTION   │   ← void
  │    VOID      │
  └──────────────┘
      |
```

SOLUTION

Jarvis

The authority layer for every action

STAGE	WHAT ACTUALLY HAPPENS
Propose	Agent or connector submits an explicit action
Approve	Human gate — authorization is deliberate
Execute	Separate step — approval is not execution
Receipt	Artifact + outcome record
Trace	Reconstruct the story end-to-end

Agents can propose anything. Authority lives in the control plane — not in the model.

PRODUCT

Product proof

Governed path in the HUD

EXECUTION CONSOLE · APPROVAL QUEUE

APPROVE	EXECUTE	RECEIPT + TRACE
Human gate before side effects	Deliberate run — not the same as approval	Artifact, log, end-to-end story

Side effects become **deliberate, attributable, and reconstructable** — what audit and procurement teams actually need to sign off.

TIMING

Why now

Three forces at once

1. **Capability** — tool-using agents are productized, not experimental
2. **Adoption** — agents touch production systems and real credentials
3. **Pressure** — governance expects **accountability**, not vibes

The missing layer: execution control with evidence.

Why this wins

> **The missing execution layer**

- Between **intent** (models, frameworks, connectors) and **side effects**
- **Authority boundaries** — who may cause what
- **Receipts + traces** for audit and ops — not log archaeology

Wedge: not the catalog — the *moment of execution*.

END STATE

Vision

> **Every production AI action is governed**

CONTROL AUDIT TRUST

- **Control** — authorization before irreversible change
- **Audit** — proof, not inference from logs
- **Trust** — humans remain the trust root for action

NEXT STEP

Thank you

Traction · team · raise · contact

Jarvis — control plane for AI execution

Demo and materials available on request