Jarvis

The control plane for AI execution

PROPOSE → APPROVE → EXECUTE → RECEIPT → TRACE

The shift

> **Agents are moving from generation to action**

- Send real outbound messages to customers
- Modify live systems and production data
- Trigger APIs, payments, and workflows

The question is not only “what did it say?” — it’s “what did it do?”

The problem

> Execution is not trusted

- Actions execute **without explicit human authorization**
- **Approval and execution collapse** into a single step
- **Logs ≠ proof** — no record of who actually authorized the action

Runtime risk is unbounded authority — not bad copy.

Market validation

> Enterprises are investing in agent infrastructure

- Massive investment in agent infrastructure (AWS, OpenAI, enterprise platforms)
- **Registries & catalogs** — what exists, who owns it
- **Orchestration & platforms** — how agents are wired
- **Governance layers** — policy, visibility, reuse

Proves the category. Does not close the execution gap.

The gap

> **Visibility ≠ control**

Platforms answer: *What can be discovered and shared?*

They do not answer:

> *What is allowed to happen — at the exact moment an agent takes action — under whose authority — with what proof?*

```
Agents & tools
      |
      ▼
Platforms (catalogs, orchestration)
      |
      ▼
  ┌──────────────┐
  │  EXECUTION   │  ← void
```

Jarvis

The authority layer for every action

Stage	What actually happens
Propose	Agent or connector submits an explicit action
Approve	Human gate — authorization is deliberate
Execute	**Separate** step — approval is not execution
Receipt	Artifact + outcome record
Trace	Reconstruct the story end-to-end

Agents can propose anything. Authority lives in the control plane.

Product proof

Governed path in the HUD

jarvis-hud · approval queue

Approve	Execute	Receipt + trace
Human gate before side effects	Deliberate run — not the same as approval	Artifact, log, end-to-end story

On video: replace with full-bleed HUD screen recording for the strongest proof; this frame reads clearly at 1080p.

Why now

Three forces at once

1. **Capability** — tool-using agents are productized, not experimental
2. **Adoption** — agents touch production systems and real credentials
3. **Pressure** — governance expects **accountability**, not vibes

The missing layer: execution control with evidence.

Why this wins

> The missing execution layer

- Between **intent** (models, frameworks, connectors) and **side effects**
- **Authority boundaries** — who may cause what
- **Receipts + traces** for audit and ops — not log archaeology

Wedge: not the catalog — the *moment of execution*.

Vision

> Every production AI action is governed

CONTROL AUDIT TRUST

- **Control** — authorization before irreversible change
- **Audit** — proof, not inference from logs
- **Trust** — humans remain the trust root for action

Thank you

Traction · team · raise · contact

Jarvis — control plane for AI execution

demo + deck available on request