Jarvis

The control plane for AI execution

PROPOSE → APPROVE → EXECUTE → RECEIPT → TRACE

The shift

> **Agents are moving from generation to action**

- Sending email and messages
- Changing systems and configuration
- Triggering workflows and integrations

The question is not only “what did it say?” — it’s “what did it do?”

The problem

> Execution is not controlled

- Side effects can run **without explicit human authorization**
- **Approval** is often conflated with **execution** — or skipped
- **Logs** record activity; they are not **proof** of who authorized what

Runtime risk is unbounded authority — not bad copy.

Market validation

> Enterprises are investing in agent infrastructure

- **Registries & catalogs** — what exists, who owns it
- **Orchestration & platforms** — how agents are wired
- **Governance layers** — policy, visibility, reuse

Proves the category. Does not close the execution gap.

The gap

> **Visibility ≠ control**

Platforms answer: *What can be discovered and shared?*

They do not answer:

> *What may happen **at the moment** an agent acts — under whose authority — with what proof?*

```
Agents & tools
      |
      ▼
Platforms (catalogs, orchestration)
      |
      ▼
┌──────────────┐
│  EXECUTION   │  ← void
```

Jarvis

Approval + proof for every action

Step	Meaning
Propose	Agent or connector submits an explicit action
Approve	Human gate — authorization is deliberate
Execute	**Separate** step — approval is not execution
Receipt	Artifact + outcome record
Trace	Reconstruct the story end-to-end

Agents can propose anything. Authority lives in the control plane.

Product proof

Governed path in the HUD

jarvis-hud · approval queue

Approve	Execute	Receipt + trace
Human gate before side effects	Deliberate run — not the same as approval	Artifact, log, end-to-end story

On video: replace with full-bleed HUD screen recording for the strongest proof; this frame reads clearly at 1080p.

Why now

Three forces at once

1. **Capability** — tool-using agents are productized, not experimental
2. **Adoption** — agents touch production systems and real credentials
3. **Pressure** — governance expects **accountability**, not vibes

The missing layer: execution control with evidence.

Why this wins

> Agent-agnostic **control plane**

- Between **intent** (models, frameworks, connectors) and **side effects**
- **Authority boundaries** — who may cause what
- **Receipts + traces** for audit and ops — not log archaeology

Wedge: not the catalog — the *moment of execution*.

Vision

> Every production AI action is governed

CONTROL AUDIT TRUST

- **Control** — authorization before irreversible change
- **Audit** — proof, not inference from logs
- **Trust** — humans remain the trust root for action

Thank you

[Traction · team · raise · contact]

Jarvis — control plane for AI execution

youtube.com / link in description